

Exercises

1. What is the time and memory complexity of the following algorithms?

```
FunctionFactors(n):  
  For i = 1 to n  
    x = n*n  
    print 'x'
```

Time: $O(n)$
Mem: $O(1)$

```
FunctionCombination2(n):  
  For i = 1 to n-1  
    For j = i to n  
      print 'i, j'
```

Time: $O(n^2)$
Mem: $O(1)$

```
FunctionCombination3(n, m):  
  For i = 1 to n  
    For j = 1 to m  
      print 'i, j'  
    For k = 1 to m  
      print 'i, k'
```

Time: $O(nm)$
Mem: $O(1)$

```
StringSearch(n, m):  
  For i = 1 to length(m) - length(n)  
    match = 0  
    For j = 1 to length(n)  
      if (n[j] = m[i+j-1])  
        match = match + 1  
    If match > length(n) * 0.95  
      print 'i'
```

Time: $O(mn)$
Mem: $O(n+m)$

```
MatchMatrix(n, m):  
  For i = 1 to length(n)  
    For j = 1 to length(m)  
      if (n[i] = m[j])  
        matrix[i, j] = 1  
      else  
        matrix[i, j] = 0
```

Time: $O(nm)$
Mem: $O(nm)$

2. Which of the following types of algorithms are inherently approximate? Which are stochastic?

approximate: greedy, randomization is often, but not always (quicksort)

stochastic: randomization, others can be stochastic but not necessarily

3. Why would you use branch and bound over exhaustive? **faster or less memory**

Today's objectives

- Evaluating strategies for short read alignment
 - String search
 - Linear
 - Hash
 - Suffix tree
 - Suffix array
 - How does short read alignment work
 - Burrows-Wheeler transform

Alignments covered & history

Classes

Pairwise vs multiple alignment

Local vs global alignment

Exhaustive vs approximate

Global, exhaustive, pairwise alignment

– Needleman S.B. and Wunsch C.D. (1970) J. Mol. Biol. 48, 443-453

Local, exhaustive, pairwise alignment

– Smith T.F. and Waterman M.S. (1981) J. Mol. Biol. 147, 195-197

Local, approximate, pairwise alignment

– BLAST: Basic Local Alignment Search Tool (Altschul et al. 1990).

Global, approximate, multiple alignment

– ClustalW: Thompson, J.D. et al. (1994) Nucleic Acids Res., 22, 4673–4680.

Local, approximate, pairwise alignment

– Bowtie/BWA: time/memory solution to short read mapping (2009/2009)

Alignments covered & history

Classes

Pairwise vs multiple alignment

Local vs global alignment

Exhaustive vs approximate

Global, exhaustive, pairwise alignment

– Needleman S.B. and Wunsch C.D. (1970) J. Mol. Biol. 48, 443-453

Local, exhaustive, pairwise alignment

– Smith T.F. and Waterman M.S. (1981) J. Mol. Biol. 147, 195-197

Local, approximate, pairwise alignment

BLAST: Basic Local Alignment Search Tool (Altschul et al. 1990).

Global, approximate, multiple alignment

ClustalW: Thompson, J.D. et al. (1994) Nucleic Acids Res., 22, 4673–4680.

Local, approximate, pairwise alignment

Bowtie/BWA: time/memory solution to short read mapping (2009/2009)

Start with
string search
-simplest
-require seeds
-find substring
-fast/memory

A simple string search

Problem: Find all substrings (x) in a string (s)

Solution: Check each position for a match

S = ACTGACTGTA

x = CTG

P1 ACT

P2 CTG

P3 GAC

P4 ACT

P5 CTG

P6 TGT

P7 GTA

A simple string search

Problem: Find all substrings (x) in a string (s)

Solution: Check each position for a match

S = ACTGACTGTA

x = CTG

P1 ACT

P2 CTG

P3 GAC

P4 ACT

P5 CTG

P6 TGT

P7 GTA

```
stringsearch( x, S)
    lx = length( x )
    ls = length( S )
    for i in 1 to ls - lx:
        if x = S[i:i+lx]
            print( i )
```

```
S = 'ACTGA'
print( S )
print( S[0:3])
```

ACTGA
ACT

S is a string; we
can retrieve slices
or indexed
positions

String vs List

```
S = 'ACTGA'
print( S )
print( S[0:3])
```

```
ACTGA
ACT
```

S is a string; we can retrieve slices or indexed positions

A is a list; we can retrieve indexed positions AND change the items

```
S[0] = 'T' # ERROR 'str' object does not support item assignment
A = list(S) # Convert string to list
A[0] = 'T' # Lists are mutable - can be changed
```

Linear search can be done with a string or list

```
print( A )
print( S )
```

```
['T', 'C', 'T', 'G', 'A']
ACTGA
```

A simple string search

Problem: Find all substrings (x) in a string (s)

Solution: Check each position for a match

S = ACTGACTGTA

x = CTG

P1 ACT

P2 CTG

P3 GAC

P4 ACT

P5 CTG

P6 TGT

P7 GTA

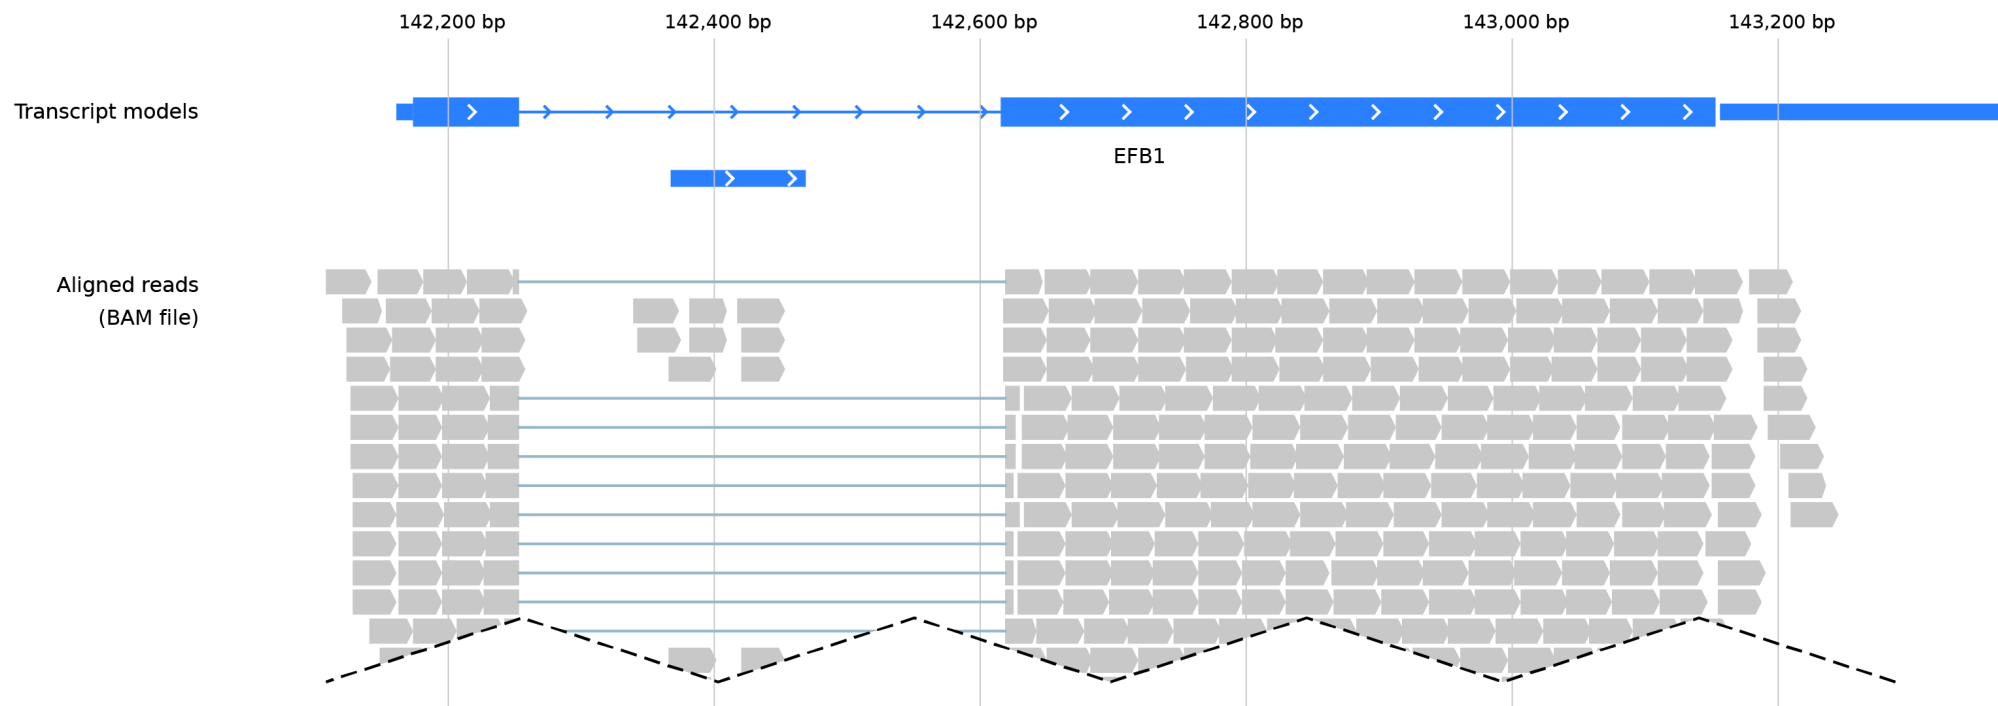
```
stringsearch( x, S)
  lx = length( x )
  ls = length( S )
  for i in 1 to ls - lx:
    if x = S[i:i+lx]
      print( i )
```

Problem --- Complexity

Time: $O(\text{length}(S))$

Memory: $O(\text{length}(S))$

Short read alignment problem



Reads

ATCGATC
TCGATCG
CAGTTGC
CGGCTCG
AATATCT

Genome

GCGCTATCGGCTCGGCTCGGATCGGCAGTTGCGTA
CGGCTCG
CGGCTCG
CAGTTGC

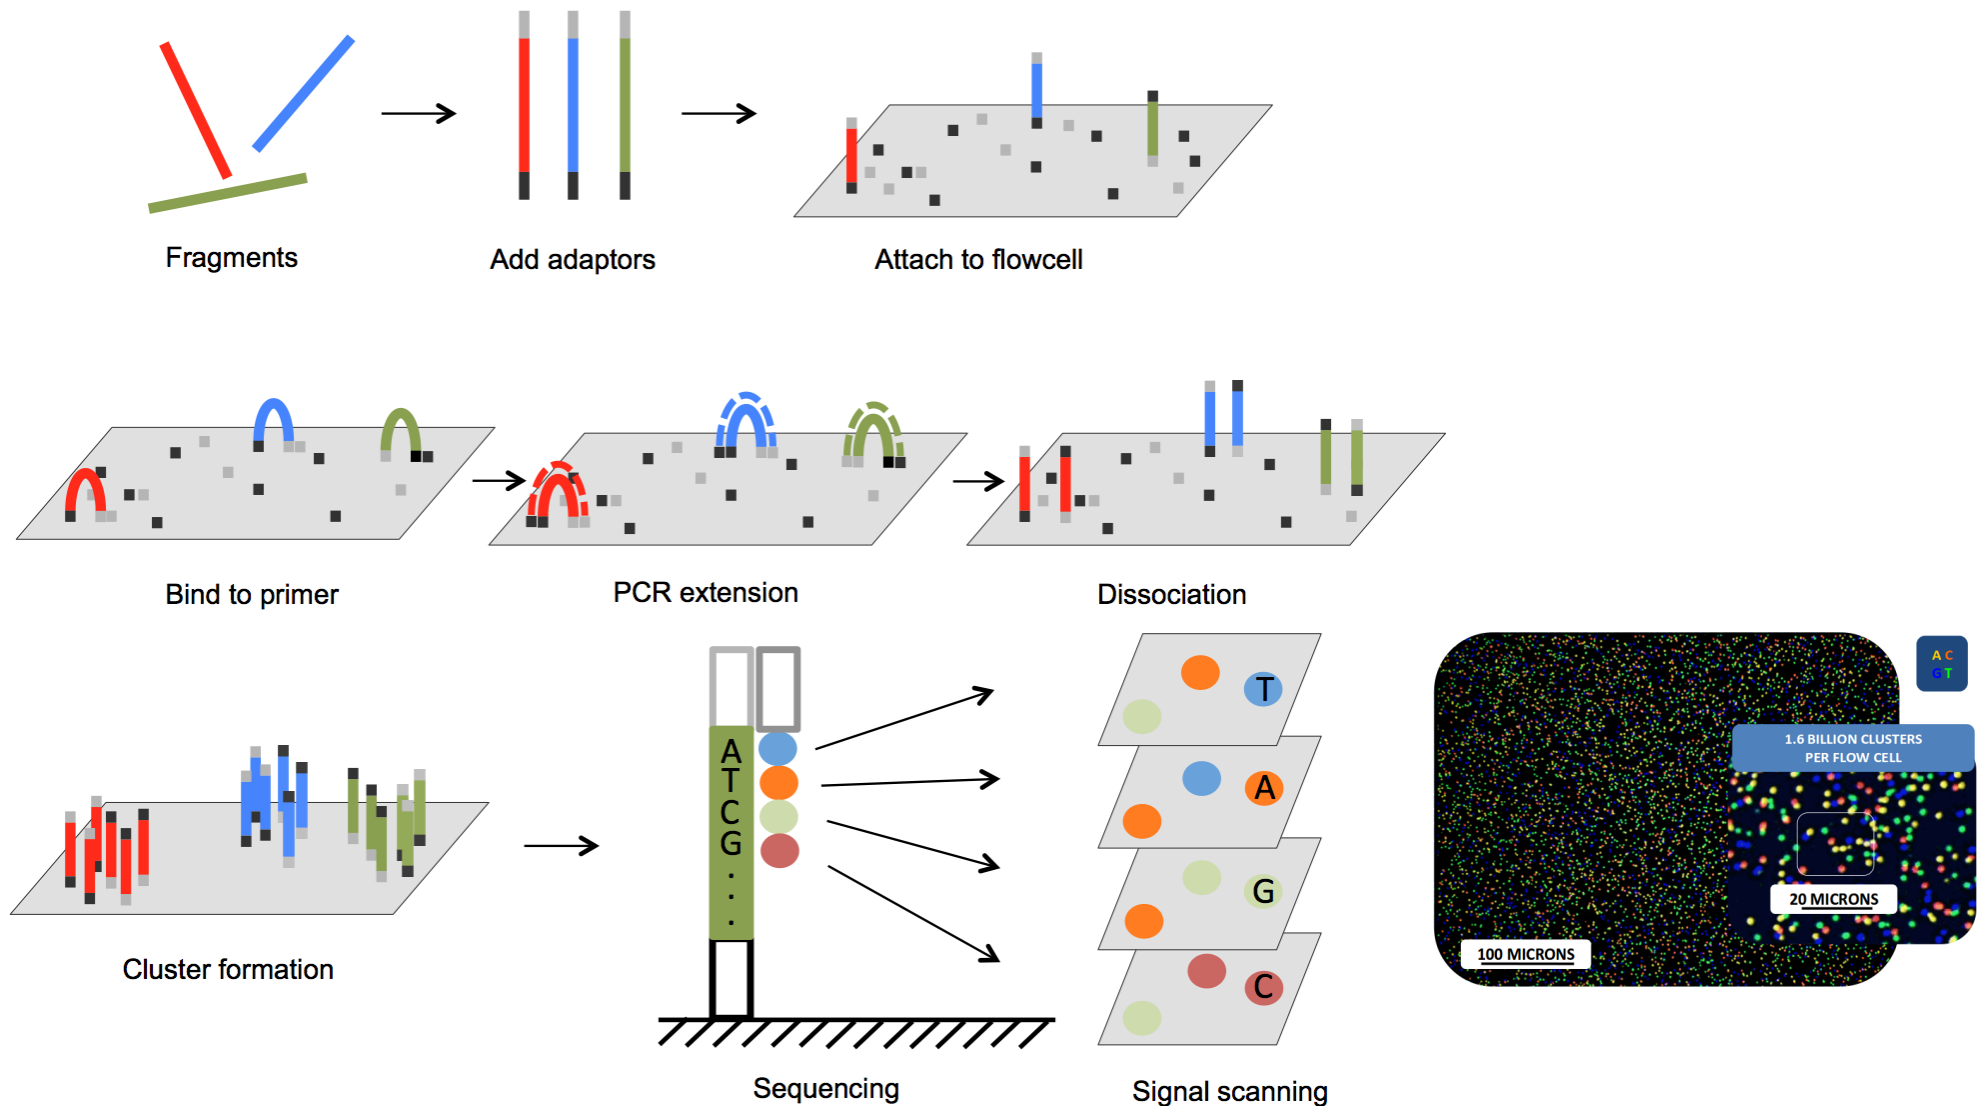


Goal: Find all exact matches of reads to genome
The problem: 10^6 reads and 10^6 to 10^9 genome size
(Reads are now 100 bp, used to be 30)
Time: $O(\text{genome size} * \text{reads}) = 10^{12}$
 $1e6/s = 10^6$ seconds or 277 hours / million reads

Reads/genome
substring/string
words/book

Illumina Sequencing

Machine	reads/run
Illumina NovaSeq	10 billion
Illumina HiSeq X	6 billion



Today's objectives

- Evaluating strategies for short read alignment

String search

- Linear
 - Hash
 - Suffix tree
 - Suffix array
- How does short read alignment work
 - Burrows-Wheeler transform

Algorithm ~ Data structure

- Linear search (array ie Python list)
- Hash (ie Python dictionary)

Algorithm History

Algorithm ~ Data structure

- Linear search (array ie Python list)
- Hash (ie Python dictionary)

Suffix Tree

- Weiner (1973)
- Knuth (father of algorithm analysis) declared it algorithm of the year

Suffix Array

- Manber and Myers (1990)
- space efficient alternative to suffix tree.

Burrows-Wheeler Transform

- Discovered by David Wheeler in 1983, published by Michael Burrows and David Wheeler in 1994
- reversible
- compression

FM-index

- Paolo Ferragina and Giovanni Manzini (2000)
- compressed full-text substring index based on the Burrows-Wheeler transform. It is used to efficiently find the number of occurrences of a pattern within the compressed text.

Short Read Alignment

Problem: Where in the genome did this short (30-100 bp) sequence come from?

- Tolerate mismatches (errors)
- Map millions of reads per hour
- Memory 2.7 GB human genome (770 Mb)
 - 1 bp = 2 bits (0/1), 1 byte = 8 bits

Solution: Burrows-Wheeler transform with FM-index

- Bowtie, BWA, SOAP2
- 25 million 35-bp reads per hour
- 1.3 GB memory

Linear search: naive, brute force

CAGATCTGC**ATGC**ATCGTAGCTAGCTACGATCGTCG

AT**G**C

ATGC

AT**G**C

ATGC

ATGC

AT**G**C

ATGC

ATGC

ATGC

ATGC

What's a faster way to
do this?

What about a hash
table?

Can we look for
inexact strings? YES

Hash Table

A **hash table** is a data structure which implements an associative array abstract data type, a structure that can map keys to values.

Python's **Dictionary** is implemented using a hash table, indexed by keys rather than numbers (List)

```
[6]: mylist = ["A", "G", "C", "T"]  
mylist[0]
```

```
[6]: 'A'
```

key

value

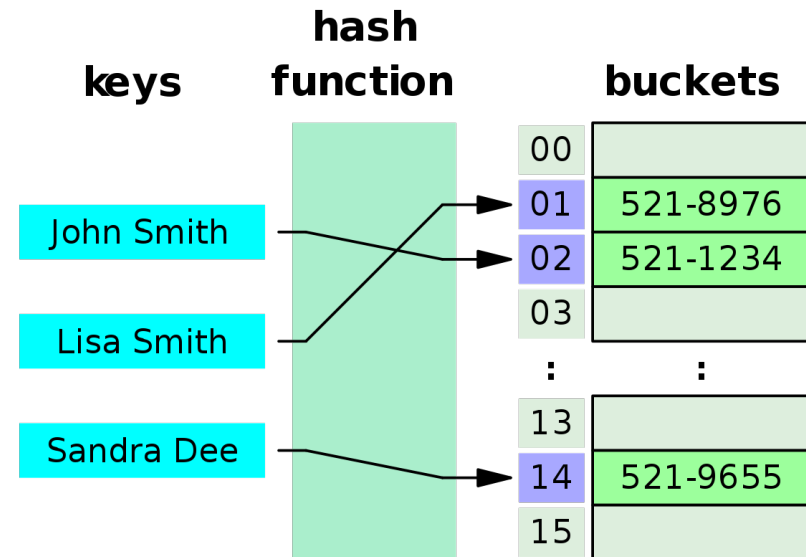
```
[7]: mydict = {"A":1, "G":2, "T":3, "C":4}  
mydict["A"]
```

```
[7]: 1
```

Hash Table

A **hash table** is a data structure which implements an associative array abstract data type, a structure that can map **keys** to **values**.

Hash table		
Type	Unordered associative array	
Invented	1953	
Time complexity in big O notation		
Algorithm	Average	Worst case
Space	$O(n)^{[1]}$	$O(n)$
Search	$O(1)$	$O(n)$



Hash collisions occur when the hash function generates the same index for more than one key. Collisions are unavoidable when hashing a large set of keys. If 2,450 keys are hashed into a million buckets, there is approximately a 95% chance of at least two of the keys being hashed to the same slot (birthday problem).

Don't worry about collisions: Python etc take care of this for you.

Hash table

S = 'AGCCCGG'

Key	Value
AG	0
GC	1
CC	2,3
CG	4
GG	5

Creating hash table

Time $O(n)$

Memory $O(n)$

Searching hash table

Time $O(1)$

Memory $O(n)$

How does complexity increase with # searches (m)?

Creating: $O(1)$

Searching: $O(m)$

Problem #1: multiple positions

Sequence	Positions
AAAAAAAAAA	32453, 64543, 76335
AAAAAAAAAAC	64534, 84323, 96536
AAAAAAAAAAG	12352, 32534, 56346
AAAAAAAAAAT	23245, 54333, 75464
AAAAAAAAACA	
AAAAAAAAACC	43523, 67543
...	
CAAAAAAAAAA	32345, 65442
CAAAAAAAAAAC	34653, 67323, 76354
...	
TCGACATGAG	54234, 67344, 75423
TCGACATGAT	11213, 22323
...	
TTTTTTTTTG	64252
TTTTTTTTTT	64246, 77355, 78453

N = length of genome

L = string length

4^L = possible strings

$N/4^L$ positions/string

If $L = 10$, $N=3e9$ or 3×10^9 :
positions/string $\sim 3,000$

Solution: Dictionary of lists

```
mydict = {"A":[1, 10, 20], "G":2, "T":3, "C":4}
mydict["A"] = [10, 20]
print(mydict["A"])
```

```
[10, 20]
```

```
print(mydict["A"][1])
```

```
20
```

Problem #2: memory

N = length of genome

L = string length

4^L = possible strings

$N/4^L$ positions/string

Humans

3.2 Gbp genome = $3e9$

If $L = 10$, $N=3e9$, ~3000 positions, slower seeding

If $L = 15$, there will be ~ 3 positions, faster seeding

But... requires 10^9 (4^{15}) hash size

1 bit = 0/1

1 byte = 8 bits

1 character = 8 bits (1 byte)

1 integer = 32 bits (4 bytes)

1 double (float) = 64 bits (8 bytes)

hash keys ($L=15$):

$10^9 * 15$ bytes = 15 Gb

hash values:

$10^9 * 4 * 3$ bytes = 12 Gb

Trade: Memory ~ Speed

1 integer * 3 positions

Problem #3: fixed string length

Suppose sequences aren't random

$L = 15$ but reads are typically 100 bp

Expected # positions = 3

But observed # positions = between 0 and 1000

Solution #1

seed and extend, but slow if 1000 positions to check

Problem #3: fixed string length

Suppose sequences aren't random

$L = 15$ but reads are typically 100 bp

Expected # positions = 3

But observed # positions = between 0 and 1000

Solution #1

seed and extend, but slow if 1000 positions to check

Read: TAGCTACGCTCTACGCTATCATGCATAAAC

Seed: TAGCTACGCT

Genome:

Seed

ATCGATTAGCTACGCTCTACGCTATCATGCATAAAC TAGCATCGCA

Extend

Problem #3: fixed string length

Suppose sequences aren't random

$L = 15$ but reads are typically 100 bp

Expected # positions = 3

But observed # positions = between 0 and 1000

Solution #1

seed and extend, but slow if 1000 positions to check

Solution #2

Longer hashes, but requires more memory and can't handle errors (string mismatches)

Data Structures

Array (Python list)
linear search, **slow**, **errors**

Time $O(n)$,
Memory $O(n)$

Hash (Python dictionary)
lookup table, **memory**, **fixed length**, **fast**

Time $O(1)$
Memory $O(n)$

Suffix tree (Mummer)

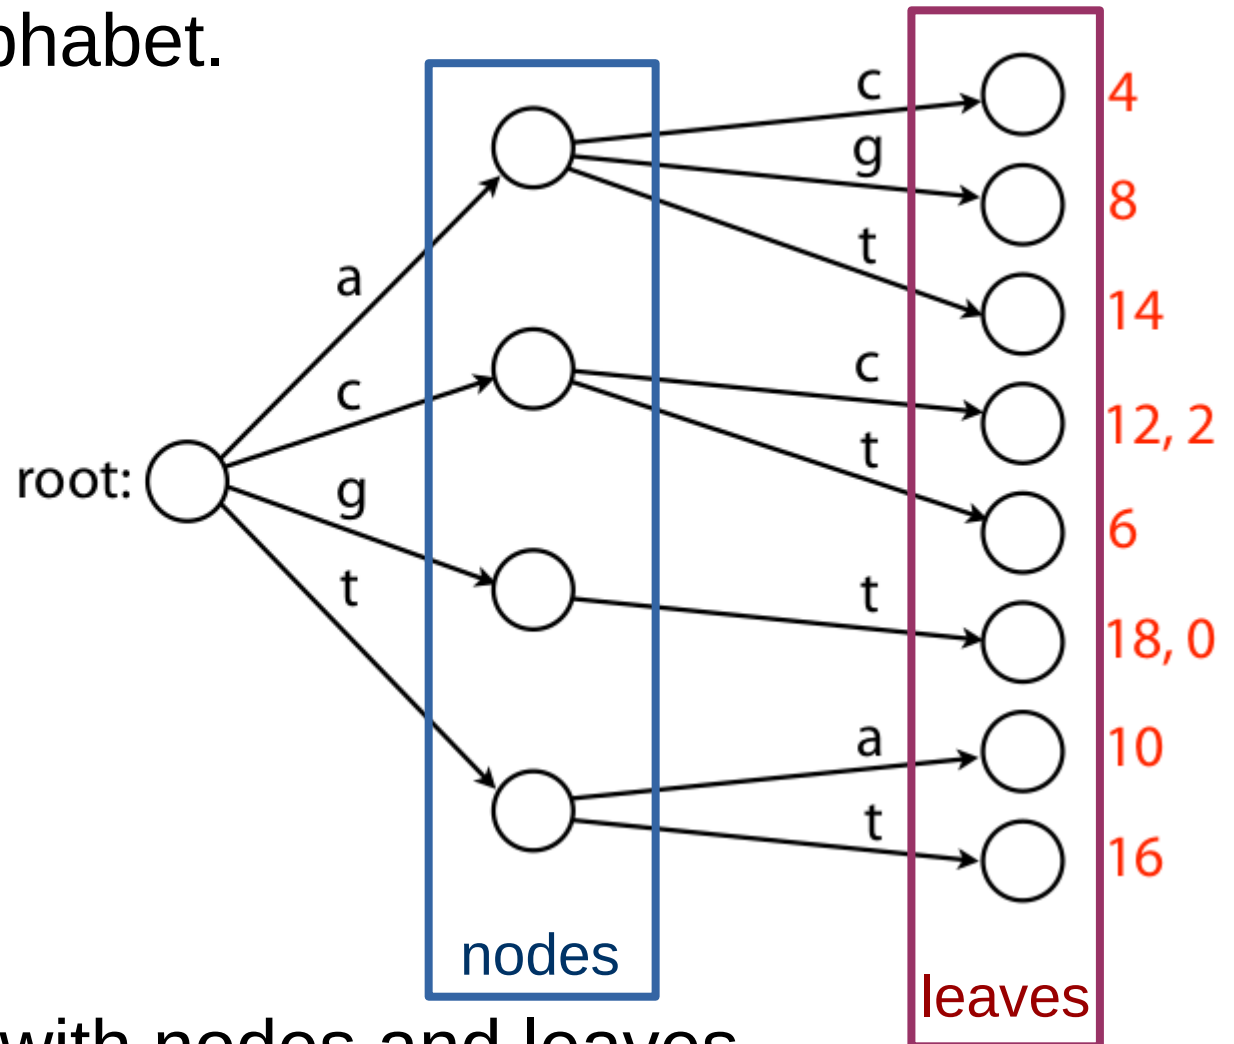
Suffix array

Burrows-Wheeler transform and FM index
(BWA and Bowtie)

Suffix Tree and Trie

Trie (radix tree or prefix tree) is a kind of search tree, an ordered tree data structure that is used to store strings over an alphabet.

ac	4
ag	8
at	14
cc	12
cc	2
ct	6
gt	18
gt	0
ta	10
tt	16



- root is empty
- keys are associated with nodes and leaves
- values are associated with leaves

Suffix trie

Suffix trie helpful to understand suffix tree

Suffix trie: build a trie containing all suffixes of text T

T: **abaaba**

abaaba

baaba

aaba

aba

ba

a



Suffixes

$n(n+1)/2$ characters

Suffix trie

Suffix trie: build a trie containing all suffixes of text T

- Add terminal character \$ to the end of T
 - $\$ < A < C < G < T$
 - ab comes before aba since $ab\$ < aba\$$
 - no suffix is a prefix for any other suffix

T: abaaba\$

abaaba\$

baaba\$

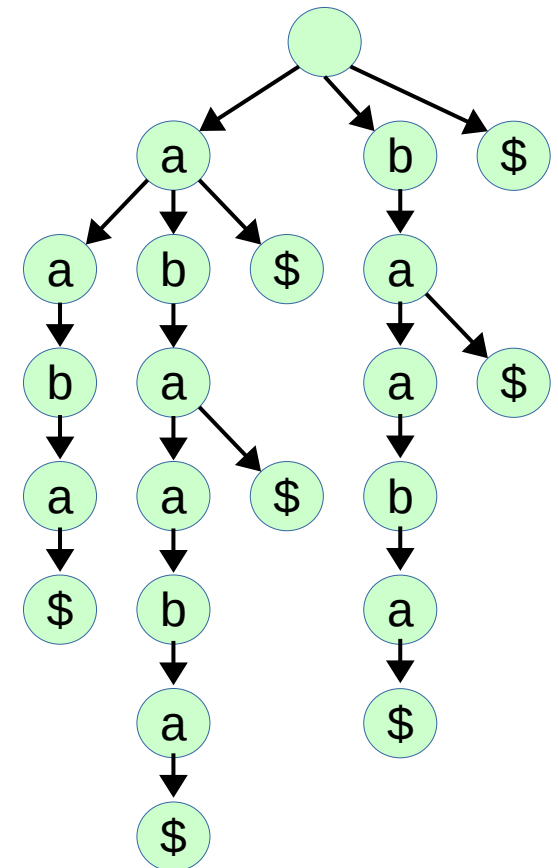
aaba\$

aba\$

ba\$

a\$

\$



Suffix trie

Suffix trie: build a trie containing all suffixes of text T

- Add terminal character \$ to the end of T
- $\$ < A < C < G < T$
- ab comes before aba since $ab\$ < aba\$$
- no suffix is a prefix for any other suffix

T: abaaba\$

abaaba\$

baaba\$

aba\$

aba\$

ba\$

a\$

\$

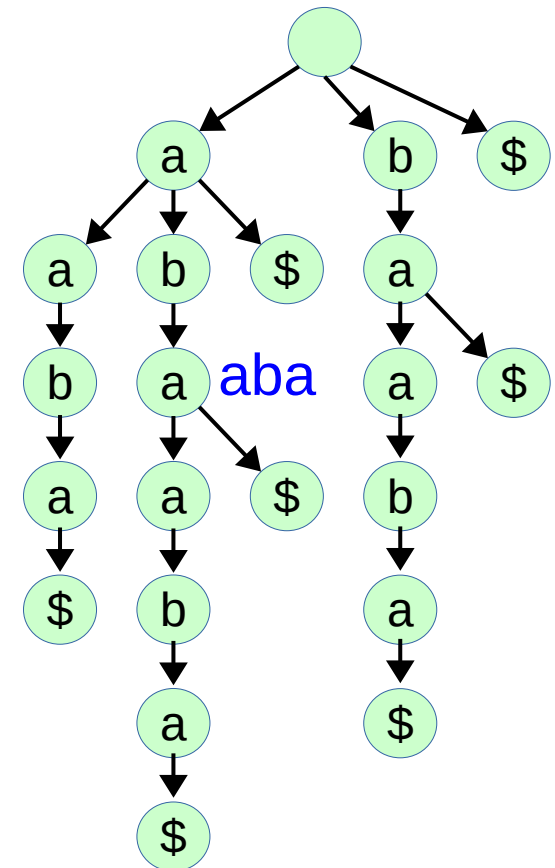
Does string 'aba' exist?

What about 'abba'?

How many times does

'aba' occur?

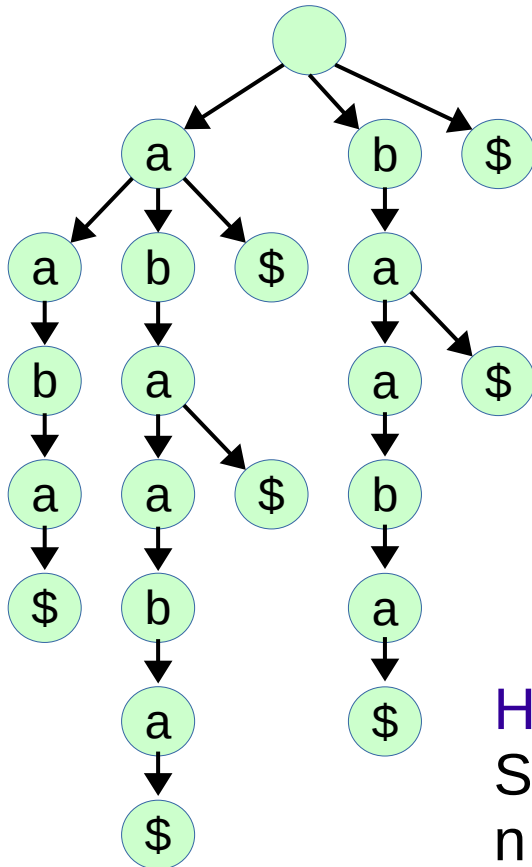
Where is 'aba'

$$n < \text{nodes} < n^2$$


Suffix Tree

- **Suffix Tree**: a compressed trie containing all the suffixes of the given text as their keys and positions in the text as their values
- Suffix trees provide linear-time solutions for the longest **common substring problem**.
- **Cost**: storing a string's suffix tree typically requires significantly more space than storing the string itself.
- Mummer: Delcher et al. (1999), Nucleic Acids Research, 27:2369-2376.

Suffix tree: compressed trie



How many nodes?

Suffix trie: $n(n+1)/2$

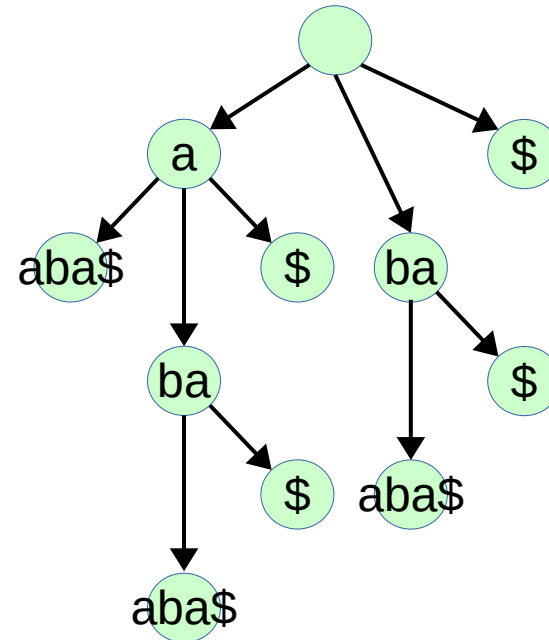
n leaf nodes

n-1 nodes (non-leaf)

$$2n-1 = O(n) \text{ (Suffix tree)}$$

But suffix tree still needs n^2 characters

keys can be stored using T and (offset, length)



Suffix tree construction

O(n) time (Ukkonen)

$O(n)$ memory

Search

$O(m)$ $m = \#$ substrings

Suffix Array

Suffix array is an array of integers specifying the lexicographic ordering of the suffixes of a string.

As only positions are stored, it is a simple, space efficient alternative to suffix trees.

Suffix	i
abaaba\$	1
baaba\$	2
aaba\$	3
aba\$	4
ba\$	5
a\$	6
\$	7

sorted	i
\$	7
a\$	6
aaba\$	3
aba\$	4
abaaba\$	1
ba\$	5
baaba\$	2

SA = [7,6,3,4,1,5,2]

Where is 'ba'?

Use binary search.

Binary Search

Binary search: finds the position of a target value within a sorted array.

Time: average $O(\log n)$

'ba' search

sorted	i
\$	7
a\$	6
aaba\$	3
aba\$	4
abaaba\$	1
ba\$	5
baaba\$	2

L



R

- Binary search compares the target value to the middle element of the array.
- If the target is greater than the **midpoint** the target cannot be to the left and so the search continues on the right
- The midpoint of the right side is found, queried against the target and this is repeated until the target is found

Binary Search

Binary search: finds the position of a target value within a sorted array.

Time: average $O(\log n)$

sorted	i
\$	7
a\$	6
aaba\$	3
aba\$	4
abaaba\$	1
ba\$	5
baaba\$	2

'ba' search

L

← L

← L

← L

R

Variations (bisect left or right)

- round the midpoint to the lower (floor) or upper (ceil) integer
- is 'ba' absent
- is 'ba' present multiple times

Data Structures

Array (Python list)
linear search, **slow**, **errors**

Time $O(n)$,
Memory $O(n)$

Hash (Python dictionary)
lookup table, **memory**, **fixed length**, **fast**

Time $O(1)$
Memory $O(n)$

Suffix tree (Mummer), **memory**, **variable length**, **fast**

Time $O(m)$
Memory $O(n)$
Construction $O(n)$

Suffix array, **memory**, **variable length**, **fast**

Time $O(m \log(n))$
Mem $O(n)$
non-Naive $O(m)$

Burrows-Wheeler transform and FM
index (BWA and Bowtie)

Speed vs Memory tradeoff
Suffix tree: a bit more
memory $\sim 20n$, but faster
than Suffix array

Time vs Memory

Brute Force (3 GB)

BANANA
BAN
ANA
NAN
ANA

Naive

Slow & Easy

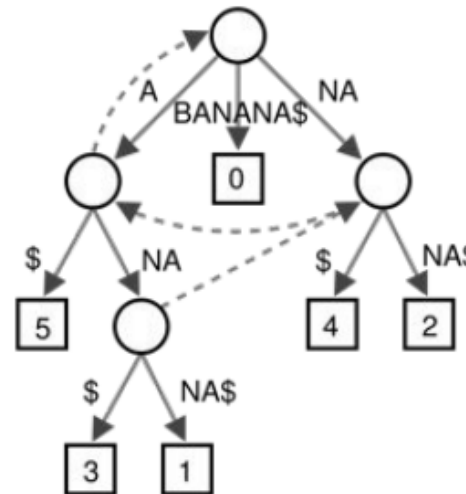
Suffix Array (>15 GB)

6	\$
5	A\$
3	ANA\$
1	ANANA\$
0	BANANA\$
4	NA\$
2	NANA\$

Vmatch, PacBio Aligner

Binary Search

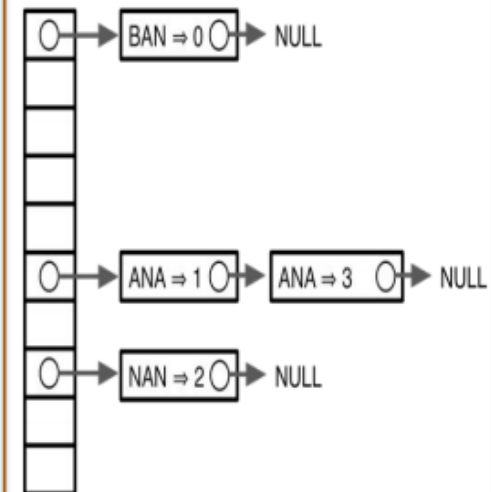
Suffix Tree (>51 GB)



MUMmer, MUMmerGPU

Tree Searching

Hash Table (>15 GB)



BLAST, MAQ, ZOOM,
RMAP, CloudBurst

Seed-and-extend

Compressed Suffix array

- Burrows-Wheeler Transform
- FM-index

Exercises

- 1) For each mapping approach:
 - i) Why not use the naive, brute force sliding window approach?
 - ii) Why not use hashes?
 - iii) Why not use suffix trees?
 - iv) Why not use suffix arrays?
- 2) Why are data structures important for the string search problem?
- 3) Which data structure would you use to search for substrings of varying lengths?
(hash, suffix trees, suffix arrays)